

Homework #2

Student: Alex Smith

Course: W261 - Machine Learning at Scale

Professor: Jimi Shanahan

Due Date: May 31

Useful resources

When completing this notebook, I found the following resources particularly useful:

- [Week 3 Slides \(https://www.dropbox.com/s/jwycz91sdi549ih/MIDS-LSML-2016-Lecture03-Map%20Reduce%20Algorithm%20Design-LiveSession3-2016-05-24.pdf?dl=0\)](https://www.dropbox.com/s/jwycz91sdi549ih/MIDS-LSML-2016-Lecture03-Map%20Reduce%20Algorithm%20Design-LiveSession3-2016-05-24.pdf?dl=0)
- [Python Sorting Dictionaries \(http://stackoverflow.com/questions/7742752/sorting-a-dictionary-by-value-then-by-key\)](http://stackoverflow.com/questions/7742752/sorting-a-dictionary-by-value-then-by-key)
- Async Lecture on Pairs & Stripes (3.9)

Libraries

The following libraries must be installed before running the below code. They can all be installed through [Pip \(https://github.com/pypa/pip\)](https://github.com/pypa/pip).

- [Scikit Learn \(http://scikit-learn.org/stable/\)](http://scikit-learn.org/stable/)
- [Numpy \(http://www.numpy.org/\)](http://www.numpy.org/)
- [Regular Expression \(https://docs.python.org/2/library/re.html\)](https://docs.python.org/2/library/re.html)
- [Pretty Table \(https://pypi.python.org/pypi/PrettyTable\)](https://pypi.python.org/pypi/PrettyTable)
- [Random \(https://docs.python.org/2/library/random.html\)](https://docs.python.org/2/library/random.html)
- [Datetime \(https://docs.python.org/2/library/datetime.html\)](https://docs.python.org/2/library/datetime.html)

Environment: Cloudera

This notebook is designed to run in a Cloudera virtual box. To set up a virtual box like this one, follow the instructions [here \(https://docs.google.com/presentation/d/1qCQM-2U2C6e584uM9kqTGr675K3_a8M1mEZaiT4Wmi8/edit#slide=id.p\)](https://docs.google.com/presentation/d/1qCQM-2U2C6e584uM9kqTGr675K3_a8M1mEZaiT4Wmi8/edit#slide=id.p). Before beginning, make sure that you have started (in the following order) from the Cloudera manager:

1. Zookeeper
2. Yarn
3. HDFS

Setting up our Hadoop file system

```
In [3]: !hdfs dfs -mkdir -p /user/cloudera/w261
```

Setting up some testing files

After you have gotten your data files into data/, let's make some testing files. This is helpful for running quick tests without churning through the whole data.

```
In [256]: !head -50 data/Consumer_Complaints.csv > data/Consumer_test.csv  
          !head -50 data/ProductPurchaseData.txt > data/ProductPurchaseData_test.txt
```

HW3.0.

How do you merge two sorted lists/arrays of records of the form [key, value]? Where is this used in Hadoop MapReduce? [Hint within the shuffle]

What is a combiner function in the context of Hadoop?

Give an example where it can be used and justify why it should be used in the context of this problem. What is the Hadoop shuffle?

After our mappers finish with the data, we feed the data into our reducers. We want to make sure that all of the values for each key are given to a single reducer. For example, in a word count program, we don't want two reducers summing values for a single word. During this phase, we can merge the two sorted lists/arrays of the form [key,value] before feeding them into the reducers. This means that we have multiple lines for a certain [key,value] pair, we can merge them by key before passing them into the reducer. This is useful because it reduces the volume of data we are moving around. We always want to reduce the volume of this data because it is risky to move data around.

To help us reduce the amount of data we move around, we can employ a combiner. The combiner is what helps merge the [key,value] pairs by the key. It is done on the mapper node before transferring the data to the reducer. We can also use a combiner on the reducer node to help simplify the data before it enters the reducer program. We can have multiple combiners at multiple points that help consolidate our data.

For example, we might want to use a combiner in the classic word count problem. We might want to combine the mapper results from the multiple mappers on a single mapper node. If we have 2 mappers running on a single node, they will both produce a list of words (the key) and a value of 1. We can combine the outputs of the two mappers by summing the values between the outputs. For instance, if mapper 1 outputs "pig 1" and mapper 2 outputs "pig 1", we can send to the reducer "pig 2". This turned two lines of data into 1.

The Hadoop **shuffle** is all the steps between mapper output and reducer input. It is helpful to break it down into five simple steps:

1. Partition, sort, combine: for a given mapper, let's sort the output and merge the [key,value] pairs by key
2. Mergesort: for all the mappers on a given node, let's sort the outputs and merge the [key,value] pairs by key
3. Send to reducer: send the data from each mapper node to the reducer
4. Mergesort again: for all the outputs from each of the mapper nodes, do another mergesort to merge the [key,value] pairs by key
5. Stream to reducer: finally send the data to the reducer program

HW3.1 consumer complaints dataset: Use Counters to do EDA (exploratory data analysis and to monitor progress)

Counters are lightweight objects in Hadoop that allow you to keep track of system progress in both the map and reduce stages of processing. By default, Hadoop defines a number of standard counters in "groups"; these show up in the jobtracker webapp, giving you information such as "Map input records", "Map output records", etc.

While processing information/data using MapReduce job, it is a challenge to monitor the progress of parallel threads running across nodes of distributed clusters. Moreover, it is also complicated to distinguish between the data that has been processed and the data which is yet to be processed. The MapReduce Framework offers a provision of user-defined Counters, which can be effectively utilized to monitor the progress of data across nodes of distributed clusters.

Use the Consumer Complaints Dataset

(https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Complaints.csv?dl=0) provide here to complete this question.

The consumer complaints dataset consists of diverse consumer complaints, which have been reported across the United States regarding various types of loans. The dataset consists of records of the form:

Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP code,Submitted via,Date received,Date sent to company,Company,Company response,Timely response?,Consumer disputed?

Here's is the first few lines of the of the Consumer Complaints Dataset:

1114245,Debt collection,Medical,Disclosure verification of debt,Not given enough info to verify debt,FL,32219,Web,11/13/2014,11/13/2014,"Choice Recovery, Inc.",Closed with explanation,Yes,
1114488,Debt collection,Medical,Disclosure verification of debt,Right to dispute notice not received,TX,75006,Web,11/13/2014,11/13/2014,"Expert Global Solutions, Inc.",In progress,Yes,
1114255,Bank account or service,Checking account,Deposits and withdrawals,,NY,11102,Web,11/13/2014,11/13/2014,"FNIS (Fidelity National Information Services, Inc.)",In progress,Yes,
1115106,Debt collection,"Other (phone, health club, etc.)",Communication tactics,Frequent or repeated calls,GA,31721,Web,11/13/2014,11/13/2014,"Expert Global Solutions, Inc.",In progress,Yes,

User-defined Counters

Now, let's use Hadoop Counters to identify the number of complaints pertaining to debt collection, mortgage and other categories (all other categories get lumped into this one) in the consumer complaints dataset. Basically produce the distribution of the Product column in this dataset using counters (limited to 3 counters here).

Hadoop offers Job Tracker, an UI tool to determine the status and statistics of all jobs. Using the job tracker UI, developers can view the Counters that have been created. Screenshot your job tracker UI as your job completes and include it here. Make sure that your user defined counters are visible.

Mapper function

We write a mapper function that takes an input of consumer complaints and outputs the category (Debt collection, Mortgage, or Other) with the number 1. This is the [key,value] pair. The mapper also writes to Hadoop counters for each of these categories.

```
In [1]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.1

# import the system library to read from
# the input and also write to stderr
import sys

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by commas since
    # this file is a CSV
    line = line.strip().split(",")

    # set the category as the second
    # item in the line
    category = line[1]

    # check which of the 3 categories
    # the line falls into and print that
    # with a 1 to the standard output
    # also print the counter to the
    # standard error so that we can double
    # check our results with the counter
    if category == "Debt collection":
        print "Debt collection\t1"
        sys.stderr.write("reporter:counter:Complaints,Debt collection,1\n")
    elif category == "Mortgage":
        print "Mortgage\t1"
        sys.stderr.write("reporter:counter:Complaints,Mortgage,1\n")
    else:
        print "Other\t1"
        sys.stderr.write("reporter:counter:Complaints,Other,1\n")
```

Writing mapper.py

Reducer function

We write a reducer function to sum across the categories, and output the number of records in each of the three categories: debt collection, mortgage, and other.

```
In [6]: %%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.1

# import the system library to read from
# the input and also write to the stderr
import sys

# create a dictionary to store the categories
# and their respective counts
categories = {}

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by tabs
    line = line.strip().split("\t")

    # set the category as the first
    # item in the line
    category = line[0]

    # set the count as the second
    # item in the line
    count = int(line[1])

    # check to see if the category already
    # exists in the dictionary
    if category not in categories:

        # if it's not, add it, and initialize
        # it with a count of 0
        categories[category] = 0

    # increment the count
    categories[category] = categories[category] + \
count

# print the outputs
for category in categories.keys():
    print category, "\t", categories[category]
```

Overwriting reducer.py

Make the files executable and run the MapReduce job in Hadoop

```
In [7]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_Complaints.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-1
!rm -r /home/cloudera/w261/Outputs/Out_3_1

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-file /home/cloudera/w261/mapper.py -mapper /home/cloudera/w261/mapper.py \
-file /home/cloudera/w261/reducer.py -reducer /home/cloudera/w261/reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-1

# copy the output file to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-1/part-00000 /home/cloudera/w261/Outputs/

# rename the file
!mv /home/cloudera/w261/Outputs/part-00000 /home/cloudera/w261/Outputs/Out_3_1
```



```

16/05/29 17:23:40 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_Complaints.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_Complaints.csv1464567820379
16/05/29 17:23:48 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-1' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-11464567828574
rm: cannot remove `/home/cloudera/w261/Outputs/Out_3_1': No such file or directory
16/05/29 17:23:50 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [/home/cloudera/w261/mapper.py, /home/cloudera/w261/reducer.py] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob5786349409546119012.jar tmpDir=null
16/05/29 17:23:51 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/05/29 17:23:52 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/05/29 17:23:53 INFO mapred.FileInputFormat: Total input paths to process : 1
16/05/29 17:23:53 INFO mapreduce.JobSubmitter: number of splits:2
16/05/29 17:23:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464566418241_0003
16/05/29 17:23:54 INFO impl.YarnClientImpl: Submitted application application_1464566418241_0003
16/05/29 17:23:54 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464566418241_0003/
16/05/29 17:23:54 INFO mapreduce.Job: Running job: job_1464566418241_0003
16/05/29 17:24:02 INFO mapreduce.Job: Job job_1464566418241_0003 running in uber mode : false
16/05/29 17:24:02 INFO mapreduce.Job:  map 0% reduce 0%
16/05/29 17:24:20 INFO mapreduce.Job:  map 67% reduce 0%
16/05/29 17:24:21 INFO mapreduce.Job:  map 100% reduce 0%
16/05/29 17:24:34 INFO mapreduce.Job:  map 100% reduce 100%
16/05/29 17:24:35 INFO mapreduce.Job: Job job_1464566418241_0003 completed successfully
16/05/29 17:24:35 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=187260
        FILE: Number of bytes written=736181
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50946999
        HDFS: Number of bytes written=54
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=407
0528
        Total time spent by all reduces in occupied slots (ms)=

```

1299712

```

Total time spent by all map tasks (ms)=31801
Total time spent by all reduce tasks (ms)=10154
Total vcore-seconds taken by all map tasks=31801
Total vcore-seconds taken by all reduce tasks=10154
Total megabyte-seconds taken by all map tasks=4070528
Total megabyte-seconds taken by all reduce tasks=129971

```

2

Map-Reduce Framework

```

Map input records=312913
Map output records=312913
Map output bytes=3324280
Map output materialized bytes=187302
Input split bytes=252
Combine input records=0
Combine output records=0
Reduce input groups=3
Reduce shuffle bytes=187302
Reduce input records=312913
Reduce output records=3
Spilled Records=625826
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=390
CPU time spent (ms)=12030
Physical memory (bytes) snapshot=390012928
Virtual memory (bytes) snapshot=2212536320
Total committed heap usage (bytes)=147324928

```

Complaints

```

Debt collection=44372
Mortgage=125752
Other=142789

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=50946747

```

File Output Format Counters

```

Bytes Written=54

```

```

16/05/29 17:24:35 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-1

```

Print the output to the notebook

```
In [9]: !cat Outputs/Out_3_1 | sort -k1,1
```

```

Debt collection      44372
Mortgage             125752
Other                142789

```

We can compare the output from the MapReduce function with what the counters calculated. We see that we get the same result!

virtual memory (bytes) snapshot		1700010120	172520132	221230020
	Name	Map	Reduce	Total
Complaints	Debt collection	44372	0	44372
	Mortgage	125752	0	125752
	Other	142789	0	142789

HW 3.2 (a) Analyze the performance of your Mappers, Combiners and Reducers using Counters

For this brief study the Input file will be one record (the next line only):

foo foo quux labs foo bar quux

Perform a word count analysis of this single record dataset using a Mapper and Reducer based WordCount (i.e., no combiners are used here) using user defined Counters to count up how many time the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing this word count job. The answer should be 1 and 4 respectively. Please explain.

Note: we split the instructions between different cells because each instruction is relatively self-contained.

Write the line to a file

```
In [13]: !echo "foo foo quux labs foo bar quux" > data/Input_3_2_a
```

Mapper function

This function takes an input, splits the line into words by white space, and outputs a key,value pair where the key is the word and the value is 1.

```
In [15]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.2(a)

# import the system library to read from
# the input and also write to stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Mapper,1\n")

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by spaces
    line = line.strip().split()

    # print out the key-value pair as a tab
    # delimited list of word and 1
    for word in line:
        print word, "\t", 1
```

Overwriting mapper.py

Reducer function

This function merges the counts for each word and outputs a list of words with their associated counts.

```
In [16]: %%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.2(a)

# import the system library to read from
# the input and also write to the stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Reducer,1\n")

# create a dictionary to store the word
# counts
words = {}

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by tabs
    line = line.strip().split("\t")

    # set the category as the first
    # item in the line
    word = line[0]

    # set the count as the second
    # item in the line
    count = int(line[1])

    # check to see if the word already
    # exists in the dictionary
    if word not in words:

        # if it's not, add it, and initialize
        # it with a count of 0
        words[word] = 0

    # increment the count
    words[word] = words[word] + count

# print the outputs
for word in words.keys():
    print word, "\t", words[word]
```

Overwriting reducer.py

Make the files executable and run the MapReduce job in Hadoop

```
In [47]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Input_3_2_a /user/clou
dera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-a
!rm -r /home/cloudera/w261/Outputs/Out_3_2-a

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-a

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-stre
aming-mr1.jar \
-D mapred.map.tasks=1 -D mapred.reduce.tasks=4 \
-file /home/cloudera/w261/mapper.py -mapper /home/cloudera/w261/mappe
r.py \
-file /home/cloudera/w261/reducer.py -reducer /home/cloudera/w261/redu
cer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-a

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-a/* /home/clouder
a/w261/Outputs/Out_3_2-a
```



```

16/05/29 19:34:20 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstar
t.cloudera:8020/user/cloudera/w261/Input_3_2_a' to trash at: hdfs://qui
ckstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/I
nput_3_2_a1464575660216
16/05/29 19:34:26 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstar
t.cloudera:8020/user/cloudera/w261-output-3-2-a' to trash at: hdfs://qu
ickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-
output-3-2-a1464575666879
16/05/29 19:34:28 WARN streaming.StreamJob: -file option is deprecated,
please use generic option -files instead.
packageJobJar: [/home/cloudera/w261/mapper.py, /home/cloudera/w261/redu
cer.py] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob4
194598760713788390.jar tmpDir=null
16/05/29 19:34:30 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/10.0.2.15:8032
16/05/29 19:34:30 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/10.0.2.15:8032
16/05/29 19:34:31 INFO mapred.FileInputFormat: Total input paths to pro
cess : 1
16/05/29 19:34:31 INFO mapreduce.JobSubmitter: number of splits:1
16/05/29 19:34:31 INFO Configuration.deprecation: mapred.reduce.tasks i
s deprecated. Instead, use mapreduce.job.reduces
16/05/29 19:34:31 INFO Configuration.deprecation: mapred.map.tasks is d
eprecated. Instead, use mapreduce.job.maps
16/05/29 19:34:31 INFO mapreduce.JobSubmitter: Submitting tokens for jo
b: job_1464566418241_0007
16/05/29 19:34:32 INFO impl.YarnClientImpl: Submitted application appli
cation_1464566418241_0007
16/05/29 19:34:32 INFO mapreduce.Job: The url to track the job: http://
quickstart.cloudera:8088/proxy/application_1464566418241_0007/
16/05/29 19:34:32 INFO mapreduce.Job: Running job: job_1464566418241_00
07
16/05/29 19:34:42 INFO mapreduce.Job: Job job_1464566418241_0007 runnin
g in uber mode : false
16/05/29 19:34:42 INFO mapreduce.Job: map 0% reduce 0%
16/05/29 19:34:57 INFO mapreduce.Job: map 100% reduce 0%
16/05/29 19:35:18 INFO mapreduce.Job: map 100% reduce 50%
16/05/29 19:35:32 INFO mapreduce.Job: map 100% reduce 100%
16/05/29 19:35:33 INFO mapreduce.Job: Job job_1464566418241_0007 comple
ted successfully
16/05/29 19:35:33 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=132
        FILE: Number of bytes written=602905
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=145
        HDFS: Number of bytes written=34
        HDFS: Number of read operations=15
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=8
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=4
        Data-local map tasks=1
        Total time spent by all maps in occupied slots (ms)=168

```

4864

Total time spent by all reduces in occupied slots (ms)=

7617536

Total time spent by all map tasks (ms)=13163

Total time spent by all reduce tasks (ms)=59512

Total vcore-seconds taken by all map tasks=13163

Total vcore-seconds taken by all reduce tasks=59512

Total megabyte-seconds taken by all map tasks=1684864

Total megabyte-seconds taken by all reduce tasks=761753

6

Map-Reduce Framework

Map input records=1

Map output records=7

Map output bytes=52

Map output materialized bytes=116

Input split bytes=114

Combine input records=0

Combine output records=0

Reduce input groups=4

Reduce shuffle bytes=116

Reduce input records=7

Reduce output records=4

Spilled Records=14

Shuffled Maps =4

Failed Shuffles=0

Merged Map outputs=4

GC time elapsed (ms)=639

CPU time spent (ms)=7150

Physical memory (bytes) snapshot=582488064

Virtual memory (bytes) snapshot=3701907456

Total committed heap usage (bytes)=232259584

MyJob

Mapper=1

Reducer=4

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Input Format Counters

Bytes Read=31

File Output Format Counters

Bytes Written=34

16/05/29 19:35:33 INFO streaming.StreamJob: Output directory: /user/clo
 udera/w261-output-3-2-a

Print the output

```
In [48]: !cat Outputs/Out_3_2-a/* | sort -n -r -k2
```

```
foo      3
quux     2
labs     1
bar      1
```

Number of counters

We use a single mapper and 4 reducers. In this simple problem, we use a single mapper because our input file is small and does not need to be chunked. We didn't need to use 4 reducers, but we go ahead with 4 so that each reducer handles 1 word and only 1 word. We might want to segregate each reducer to a single word if we were analyzing a very large corpus. This would be an intuitive way to break up the output files.

	Name ^	Map ↕	Reduce ↕	Total ↕
MyJob	Mapper	1	0	1
	Reducer	0	4	4

HW 3.2 (b)

Please use multiple mappers and reducers for these jobs (at least 2 mappers and 2 reducers). Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper and Reducer based WordCount (i.e., no combiners used anywhere) using user defined Counters to count up how many times the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job?

Mapper function

This function takes the string from the issue column and outputs a line for each word following by the integer 1.

```
In [259]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.2(b)

# import the system library to read from
# the input and also write to stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Mapper,1\n")

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by spaces
    line = line.strip().split(" ")

    # grab the words we're interested in
    words = line[3].split()

    # print out the key-value pair as a tab
    # delimited list of word and 1
    for word in words:
        print word, "\t", 1
```

Overwriting mapper.py

Reducer function

This function takes the inputs of key,value pairs of words and counts and merges the counts by word to generate a list of total counts for each word.

```
In [260]: %%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.2(b)

# import the system library to read from
# the input and also write to the stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Reducer,1\n")

# create a dictionary to store the word
# counts
words = {}

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by tabs
    line = line.strip().split("\t")

    # set the category as the first
    # item in the line
    word = line[0]

    # set the count as the second
    # item in the line
    count = int(line[1])

    # check to see if the word already
    # exists in the dictionary
    if word not in words:

        # if it's not, add it, and initialize
        # it with a count of 0
        words[word] = 0

    # increment the count
    words[word] = words[word] + count

# print the outputs
for word in words.keys():
    print word, "\t", words[word]
```

Overwriting reducer.py

Make the files executable and run the test data in Hadoop

```
In [10]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_test.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-b
!rm -r /home/cloudera/w261/Outputs/Out_3_2-b

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-b

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=2 \
-file /home/cloudera/w261/mapper.py -mapper /home/cloudera/w261/mapper.py \
-file /home/cloudera/w261/reducer.py -reducer /home/cloudera/w261/reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-b

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-b/* /home/cloudera/w261/Outputs/Out_3_2-b
```



```
16/05/30 12:22:26 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_test.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_test.csv
16/05/30 12:22:40 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-2-b' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-2-b
16/05/30 12:22:42 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [/home/cloudera/w261/mapper.py, /home/cloudera/w261/reducer.py] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob3598338797964190338.jar tmpDir=null
16/05/30 12:22:45 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 12:22:45 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 12:22:47 INFO mapred.FileInputFormat: Total input paths to process : 1
16/05/30 12:22:48 INFO mapreduce.JobSubmitter: number of splits:2
16/05/30 12:22:48 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
16/05/30 12:22:48 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/05/30 12:22:49 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0002
16/05/30 12:22:49 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0002
16/05/30 12:22:50 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0002/
16/05/30 12:22:50 INFO mapreduce.Job: Running job: job_1464634906532_0002
16/05/30 12:23:17 INFO mapreduce.Job: Job job_1464634906532_0002 running in uber mode : false
16/05/30 12:23:17 INFO mapreduce.Job:  map 0% reduce 0%
16/05/30 12:24:12 INFO mapreduce.Job:  map 100% reduce 0%
16/05/30 12:24:33 INFO mapreduce.Job:  map 100% reduce 100%
16/05/30 12:24:36 INFO mapreduce.Job: Job job_1464634906532_0002 completed successfully
16/05/30 12:24:36 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=1768
        FILE: Number of bytes written=486360
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=61947
        HDFS: Number of bytes written=1195
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=136
```

32256


```

Total time spent by all reduces in occupied slots (ms)=
4535424
Total time spent by all map tasks (ms)=106502
Total time spent by all reduce tasks (ms)=35433
Total vcore-seconds taken by all map tasks=106502
Total vcore-seconds taken by all reduce tasks=35433
Total megabyte-seconds taken by all map tasks=13632256
Total megabyte-seconds taken by all reduce tasks=453542

```

4

Map-Reduce Framework

```

Map input records=249
Map output records=1083
Map output bytes=11212
Map output materialized bytes=2410
Input split bytes=240
Combine input records=0
Combine output records=0
Reduce input groups=102
Reduce shuffle bytes=2410
Reduce input records=1083
Reduce output records=102
Spilled Records=2166
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=2022
CPU time spent (ms)=5850
Physical memory (bytes) snapshot=491601920
Virtual memory (bytes) snapshot=2949853184
Total committed heap usage (bytes)=190316544

```

MyJob

```

Mapper=2
Reducer=2

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=61707

```

File Output Format Counters

```

Bytes Written=1195

```

```

16/05/30 12:24:36 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-b

```

Run the full data through Hadoop

```
In [261]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_Complaints.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-b
!rm -r /home/cloudera/w261/Outputs/Out_3_2-b

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-b

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=2 \
-file /home/cloudera/w261/mapper.py -mapper /home/cloudera/w261/mapper.py \
-file /home/cloudera/w261/reducer.py -reducer /home/cloudera/w261/reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-b

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-b/* /home/cloudera/w261/Outputs/Out_3_2-b
```



```
16/06/04 09:56:24 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_Complaints.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_Complaints.csv
16/06/04 09:56:35 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-2-b' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-2-b1465059395441
16/06/04 09:56:37 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [/home/cloudera/w261/mapper.py, /home/cloudera/w261/reducer.py] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob1815764423982164672.jar tmpDir=null
16/06/04 09:56:39 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 09:56:40 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 09:56:42 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 09:56:42 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 09:56:42 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
16/06/04 09:56:42 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/06/04 09:56:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0019
16/06/04 09:56:44 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0019
16/06/04 09:56:44 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0019/
16/06/04 09:56:44 INFO mapreduce.Job: Running job: job_1464634906532_0019
16/06/04 09:57:04 INFO mapreduce.Job: Job job_1464634906532_0019 running in uber mode : false
16/06/04 09:57:04 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 09:57:58 INFO mapreduce.Job:  map 64% reduce 0%
16/06/04 09:58:02 INFO mapreduce.Job:  map 67% reduce 0%
16/06/04 09:58:04 INFO mapreduce.Job:  map 83% reduce 0%
16/06/04 09:58:05 INFO mapreduce.Job:  map 92% reduce 0%
16/06/04 09:58:07 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 09:58:31 INFO mapreduce.Job:  map 100% reduce 84%
16/06/04 09:58:34 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 09:58:41 INFO mapreduce.Job: Job job_1464634906532_0019 completed successfully
16/06/04 09:58:41 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=916715
        FILE: Number of bytes written=1976599
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50946999
        HDFS: Number of bytes written=2659
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
```

```

Launched map tasks=2
Launched reduce tasks=2
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)=139
36128
Total time spent by all reduces in occupied slots (ms)=
7184640
Total time spent by all map tasks (ms)=108876
Total time spent by all reduce tasks (ms)=56130
Total vcore-seconds taken by all map tasks=108876
Total vcore-seconds taken by all reduce tasks=56130
Total megabyte-seconds taken by all map tasks=13936128
Total megabyte-seconds taken by all reduce tasks=718464
0

```

Map-Reduce Framework

```

Map input records=312913
Map output records=978634
Map output bytes=10407527
Map output materialized bytes=590545
Input split bytes=252
Combine input records=0
Combine output records=0
Reduce input groups=182
Reduce shuffle bytes=590545
Reduce input records=978634
Reduce output records=182
Spilled Records=2483605
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=2924
CPU time spent (ms)=20260
Physical memory (bytes) snapshot=497647616
Virtual memory (bytes) snapshot=2973028352
Total committed heap usage (bytes)=188743680

```

MyJob

```

Mapper=2
Reducer=2

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=50946747

```

File Output Format Counters

```

Bytes Written=2659

```

```

16/06/04 09:58:41 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-b

```

Sample the output file

Let's print a sample from one of the output files

```
In [263]: !head Outputs/Out_3_2-b/part-00000
```

```
"Account      16205
the          6248
increase/decrease  1149
APR          3431
Account      350
promised     274
Overlimit    127
process      5505
issue       1098
Debt         1343
```

Mapper and reducer counters

We used a total of 2 reducers and 2 mappers. We can use 2 reducers easily because we are not looking at connections between words, only the count of each word. Because Hadoop sorts by keys, each reducer will get all the records for a single word.

MyJob	Name	Map	Reduce	Total
	Mapper	2	0	2
	Reducer	0	2	2

3.2 (c)

Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper, Reducer, and standalone combiner (i.e., not an in-memory combiner) based WordCount using user defined Counters to count up how many time the mapper, combiner, reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

Mapper function

This function takes the words in the issue column and outputs each word with an integer 1.

```
In [264]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.2(c)

# import the system library to read from
# the input and also write to stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Mapper,1\n")

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by spaces
    line = line.strip().split(" ")

    # grab the words we're interested in
    words = line[3].split()

    # print out the key-value pair as a tab
    # delimited list of word and 1
    for word in words:
        print word, "\t", 1
```

Overwriting mapper.py

Reducer function

This function takes the words outputted from the mapper and merges based on the word to get a sum of counts for each word.

```
In [265]: %%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.2(c)

# import the system library to read from
# the input and also write to the stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Reducer,1\n")

# create a dictionary to store the word
# counts
words = {}

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by tabs
    line = line.strip().split("\t")

    # set the category as the first
    # item in the line
    word = line[0]

    # set the count as the second
    # item in the line
    count = int(line[1])

    # check to see if the word already
    # exists in the dictionary
    if word not in words:

        # if it's not, add it, and initialize
        # it with a count of 0
        words[word] = 0

    # increment the count
    words[word] = words[word] + count

# print the outputs
for word in words.keys():
    print word, "\t", words[word]
```

Overwriting `reducer.py`

Make executable and run the test data in hadoop

We add our reducer as a stand-alone combiner.

```
In [15]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_test.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-c
!rm -r /home/cloudera/w261/Outputs/Out_3_2-c

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-c

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=2 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-combiner reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-c

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-c/* /home/cloudera/w261/Outputs/Out_3_2-c
```



```

16/05/30 12:54:01 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_test.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_test.csv1464638041156
16/05/30 12:54:09 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-2-c' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-2-c
16/05/30 12:54:13 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/05/30 12:54:13 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob247530583643913124.jar tmpDir=null
16/05/30 12:54:14 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 12:54:14 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 12:54:16 INFO mapred.FileInputFormat: Total input paths to process : 1
16/05/30 12:54:16 INFO mapreduce.JobSubmitter: number of splits:2
16/05/30 12:54:17 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0004
16/05/30 12:54:17 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0004
16/05/30 12:54:17 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0004/
16/05/30 12:54:17 INFO mapreduce.Job: Running job: job_1464634906532_0004
16/05/30 12:54:26 INFO mapreduce.Job: Job job_1464634906532_0004 running in uber mode : false
16/05/30 12:54:26 INFO mapreduce.Job:  map 0% reduce 0%
16/05/30 12:54:52 INFO mapreduce.Job:  map 100% reduce 0%
16/05/30 12:55:14 INFO mapreduce.Job:  map 100% reduce 100%
16/05/30 12:55:16 INFO mapreduce.Job: Job job_1464634906532_0004 completed successfully
16/05/30 12:55:16 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=1476
        FILE: Number of bytes written=486790
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=61947
        HDFS: Number of bytes written=1297
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=5956480
        Total time spent by all reduces in occupied slots (ms)=4658816
        Total time spent by all map tasks (ms)=46535

```

```

Total time spent by all reduce tasks (ms)=36397
Total vcore-seconds taken by all map tasks=46535
Total vcore-seconds taken by all reduce tasks=36397
Total megabyte-seconds taken by all map tasks=5956480
Total megabyte-seconds taken by all reduce tasks=465881

```

6

Map-Reduce Framework

```

Map input records=249
Map output records=1083
Map output bytes=11212
Map output materialized bytes=1940
Input split bytes=240
Combine input records=1083
Combine output records=159
Reduce input groups=156
Reduce shuffle bytes=1940
Reduce input records=159
Reduce output records=102
Spilled Records=318
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=667
CPU time spent (ms)=4750
Physical memory (bytes) snapshot=488947712
Virtual memory (bytes) snapshot=2959560704
Total committed heap usage (bytes)=191889408

```

MyJob

```

Mapper=2
Reducer=6

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=61707

```

File Output Format Counters

```

Bytes Written=1297

```

```

16/05/30 12:55:16 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-c

```

Run the full data through Hadoop

```
In [266]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_Complaints.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-c
!rm -r /home/cloudera/w261/Outputs/Out_3_2-c

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-c

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=2 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-combiner reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-c

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-c/* /home/cloudera/w261/Outputs/Out_3_2-c
```



```

16/06/04 10:10:41 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_Complaints.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_Complaints.csv1465060241812
16/06/04 10:10:57 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-2-c' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-2-c
16/06/04 10:11:01 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/06/04 10:11:01 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob438746980786419372.jar tmpDir=null
16/06/04 10:11:02 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:11:02 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:11:04 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 10:11:04 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 10:11:04 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0020
16/06/04 10:11:05 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0020
16/06/04 10:11:05 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0020/
16/06/04 10:11:05 INFO mapreduce.Job: Running job: job_1464634906532_0020
16/06/04 10:11:16 INFO mapreduce.Job: Job job_1464634906532_0020 running in uber mode : false
16/06/04 10:11:16 INFO mapreduce.Job: map 0% reduce 0%
16/06/04 10:11:54 INFO mapreduce.Job: map 67% reduce 0%
16/06/04 10:11:58 INFO mapreduce.Job: map 100% reduce 0%
16/06/04 10:12:18 INFO mapreduce.Job: map 100% reduce 50%
16/06/04 10:12:19 INFO mapreduce.Job: map 100% reduce 100%
16/06/04 10:12:21 INFO mapreduce.Job: Job job_1464634906532_0020 completed successfully
16/06/04 10:12:21 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=12550
        FILE: Number of bytes written=499602
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50946999
        HDFS: Number of bytes written=2841
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=10056576
        Total time spent by all reduces in occupied slots (ms)=

```


4282624

```

Total time spent by all map tasks (ms)=78567
Total time spent by all reduce tasks (ms)=33458
Total vcore-seconds taken by all map tasks=78567
Total vcore-seconds taken by all reduce tasks=33458
Total megabyte-seconds taken by all map tasks=10056576
Total megabyte-seconds taken by all reduce tasks=428262

```

4

Map-Reduce Framework

```

Map input records=312913
Map output records=978634
Map output bytes=10407527
Map output materialized bytes=6208
Input split bytes=252
Combine input records=978634
Combine output records=508
Reduce input groups=506
Reduce shuffle bytes=6208
Reduce input records=508
Reduce output records=182
Spilled Records=1369
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=1229
CPU time spent (ms)=9250
Physical memory (bytes) snapshot=485023744
Virtual memory (bytes) snapshot=2953555968
Total committed heap usage (bytes)=192937984

```

MyJob

```

Mapper=2
Reducer=8

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=50946747

```

File Output Format Counters

```

Bytes Written=2841

```

```

16/06/04 10:12:21 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-c

```

Mapper and reducer counters

We use a total of 8 counters, 6 reducer counters and 2 mapper counters. On the mapper side, we use 2 mapper counters and 4 reducer counters. On the reducer side, we use 2 reducer counters. We only have 2 mappers and 2 reducers but the counters show more because we are using the reducers as combiners. We are running the reducers as combiners twice for each mapper. In the mapper, this could be happening once in the circular buffer and once after all the output has been spilled to disk.

HW3.2 (d)

Using a single reducer: What are the top 50 most frequent terms in your word count analysis? Present the top 50 terms and their frequency and their relative frequency. Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order. Present bottom 10 tokens (least frequent items).

Mapper function

This function takes the words from the complaints data and outputs each word with an integer 1.

```
In [276]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.2(d)

# import the system library to read from
# the input and also write to stderr
import sys

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Mapper,1\n")

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by spaces
    line = line.strip().split(" ")

    # grab the words we're interested in
    words = line[3].split()

    # print out the key-value pair as a tab
    # delimited list of word and 1
    for word in words:
        print word, "\t", 1
```

Overwriting mapper.py

Reducer function

This function merges the counts for each word, by word to produce a sum of word counts. It returns a sorted list by the word count.

```
In [277]: %%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.2(d)

# import the system library to read from
# the input and also write to the stderr
# import the operator library to help sort
# the dictionary
import sys
import operator

# add a counter line for each time the
# mapper is called
sys.stderr.write("reporter:counter:MyJob,Reducer,1\n")

# create a dictionary to store the word
# counts
words = {}

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by tabs
    line = line.strip().split("\t")

    # set the word as the first
    # item in the line
    word = line[0].strip()

    # set the count as the second
    # item in the line
    count = int(line[1])

    # check to see if the word already
    # exists in the dictionary
    if word not in words:

        # if it's not, add it, and initialize
        # it with a count of 0
        words[word] = 0

    # increment the count
    words[word] = words[word] + count

# print the outputs
for word in sorted(words.items(),\
                    key=operator.itemgetter(1)):
    print word[0], "\t", word[1]
```

Overwriting reducer.py

Make the files executable and run the hadoop

```
In [20]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_test.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-d
!rm -r /home/cloudera/w261/Outputs/Out_3_2-d

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-d

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-combiner reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-d

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-d/part-00000 /home/cloudera/w261/Outputs/Out_3_2-d
```



```

16/05/30 13:28:01 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_test.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_test.csv1464640081215
rm: `/user/cloudera/w261-output-3-2-d': No such file or directory
rm: cannot remove `/home/cloudera/w261/Outputs/Out_3_2-d': No such file or directory
16/05/30 13:28:19 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/05/30 13:28:19 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob6504188710895385985.jar tmpDir=null
16/05/30 13:28:20 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 13:28:21 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 13:28:23 INFO mapred.FileInputFormat: Total input paths to process : 1
16/05/30 13:28:24 INFO mapreduce.JobSubmitter: number of splits:2
16/05/30 13:28:24 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0005
16/05/30 13:28:25 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0005
16/05/30 13:28:25 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0005/
16/05/30 13:28:25 INFO mapreduce.Job: Running job: job_1464634906532_0005
16/05/30 13:28:47 INFO mapreduce.Job: Job job_1464634906532_0005 running in uber mode : false
16/05/30 13:28:47 INFO mapreduce.Job:  map 0% reduce 0%
16/05/30 13:29:44 INFO mapreduce.Job:  map 100% reduce 0%
16/05/30 13:29:55 INFO mapreduce.Job:  map 100% reduce 100%
16/05/30 13:29:56 INFO mapreduce.Job: Job job_1464634906532_0005 completed successfully
16/05/30 13:29:56 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=1357
        FILE: Number of bytes written=365632
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=61947
        HDFS: Number of bytes written=1093
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=13684992
        Total time spent by all reduces in occupied slots (ms)=1063296
        Total time spent by all map tasks (ms)=106914
        Total time spent by all reduce tasks (ms)=8307

```

```

Total vcore-seconds taken by all map tasks=106914
Total vcore-seconds taken by all reduce tasks=8307
Total megabyte-seconds taken by all map tasks=13684992
Total megabyte-seconds taken by all reduce tasks=106329

```

6

Map-Reduce Framework

```

Map input records=249
Map output records=1083
Map output bytes=11212
Map output materialized bytes=1750
Input split bytes=240
Combine input records=1083
Combine output records=159
Reduce input groups=159
Reduce shuffle bytes=1750
Reduce input records=159
Reduce output records=102
Spilled Records=318
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=537
CPU time spent (ms)=3430
Physical memory (bytes) snapshot=389562368
Virtual memory (bytes) snapshot=2207375360
Total committed heap usage (bytes)=143130624

```

MyJob

```

Mapper=2
Reducer=3

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=61707

```

File Output Format Counters

```

Bytes Written=1093

```

```

16/05/30 13:29:56 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-d

```

Run the full data through Hadoop


```
In [278]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_Complaints.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-d
!rm -r /home/cloudera/w261/Outputs/Out_3_2-d

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-d

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-combiner reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-d

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-d/part-00000 /home/cloudera/w261/Outputs/Out_3_2-d
```



```

16/06/04 10:22:53 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_Complaints.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_Complaints.csv1465060973395
16/06/04 10:23:02 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-2-d' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-2-d
16/06/04 10:23:05 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/06/04 10:23:05 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob7367082926943165624.jar tmpDir=null
16/06/04 10:23:06 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:23:06 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:23:08 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 10:23:08 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 10:23:08 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0021
16/06/04 10:23:09 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0021
16/06/04 10:23:09 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0021/
16/06/04 10:23:09 INFO mapreduce.Job: Running job: job_1464634906532_0021
16/06/04 10:23:19 INFO mapreduce.Job: Job job_1464634906532_0021 running in uber mode : false
16/06/04 10:23:19 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 10:23:56 INFO mapreduce.Job:  map 67% reduce 0%
16/06/04 10:23:59 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 10:24:10 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 10:24:10 INFO mapreduce.Job: Job job_1464634906532_0021 completed successfully
16/06/04 10:24:10 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=9242
        FILE: Number of bytes written=377430
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50946999
        HDFS: Number of bytes written=2477
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=9445888
        Total time spent by all reduces in occupied slots (ms)=1024768

```

```

Total time spent by all map tasks (ms)=73796
Total time spent by all reduce tasks (ms)=8006
Total vcore-seconds taken by all map tasks=73796
Total vcore-seconds taken by all reduce tasks=8006
Total megabyte-seconds taken by all map tasks=9445888
Total megabyte-seconds taken by all reduce tasks=102476

```

8

Map-Reduce Framework

```

Map input records=312913
Map output records=978634
Map output bytes=10407527
Map output materialized bytes=5663
Input split bytes=252
Combine input records=978634
Combine output records=508
Reduce input groups=508
Reduce shuffle bytes=5663
Reduce input records=508
Reduce output records=182
Spilled Records=1369
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=899
CPU time spent (ms)=7400
Physical memory (bytes) snapshot=376905728
Virtual memory (bytes) snapshot=2190626816
Total committed heap usage (bytes)=145227776

```

MyJob

```

Mapper=2
Reducer=4

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=50946747

```

File Output Format Counters

```

Bytes Written=2477

```

```

16/06/04 10:24:10 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-d

```

Top 50 tokens

These 50 words appeared most frequently.

```
In [281]: !tail -50 Outputs/Out_3_2-d/part-00000 | sort -k2nr,2 -k1
```

```
"Loan      107254
modification    70487
servicing       36767
credit  36126
report  30546
Incorrect       29069
information     29069
on      29069
or      22533
debt    17966
and     16448
"Account       16205
opening        16205
Credit  14768
club      12545
health   12545
/         12386
not      12353
loan     12237
attempts     11848
collect      11848
Cont'd   11848
owed      11848
of        10885
my        10731
Deposits     10555
withdrawals  10555
Problems     9484
"Application  8625
to          8401
Billing      8158
Other       7886
disputes    6938
Communication 6920
tactics     6920
reporting   6559
lease       6337
the         6248
being      5663
by         5663
caused     5663
funds      5663
low        5663
process     5505
Disclosure   5214
verification 5214
Managing     5006
company's    4858
investigation 4858
card        4405
```

Bottom 10 tokens

These 10 words appear least frequently.

```
In [285]: !head Outputs/Out_3_2-d/part-00000 | sort -k2nr,2 -k1
```

```
credited      92
Payment       92
checks 75
Convenience   75
amt          71
day          71
wrong 71
disclosures   64
Incorrect/missing 64
Issue 1
```

HW 3.2.1

Using 2 reducers: What are the top 50 most frequent terms in your word count analysis? Present the top 50 terms and their frequency and their relative frequency. Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order. Present bottom 10 tokens (least frequent items). Please use a combiner.

Mapper function

This function takes an input of words and outputs each word with it's associated count.

```
In [286]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.2.1

# import the system library to read from
# the input and also write to stderr
import sys

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by spaces
    line = line.strip().split(",")

    # grab the words we're interested in
    words = line[3].split()

    # print out the key-value pair as a tab
    # delimited list of word and 1
    for word in words:
        print word, "\t", 1
```

Overwriting mapper.py

Reducer function

This function sums across the counts for a single word and outputs a list of word with the final counts.

```
In [287]: %%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.2.1

# import the system library to read from
# the input and also write to the stderr
# import the operator library to help sort
# the dictionary
import sys
import operator

# create a dictionary to store the word
# counts
words = {}

# loop through each line
for line in sys.stdin:

    # strip off any extra spaces and
    # split the line by tabs
    line = line.strip().split("\t")

    # set the word as the first
    # item in the line
    word = line[0].strip()

    # set the count as the second
    # item in the line
    count = int(line[1])

    # check to see if the word already
    # exists in the dictionary
    if word not in words:

        # if it's not, add it, and initialize
        # it with a count of 0
        words[word] = 0

    # increment the count
    words[word] = words[word] + count

# print the outputs
for word in sorted(words.items(),\
                    key=operator.itemgetter(1)):
    print word[0], "\t", word[1]
```

Overwriting reducer.py

Test the code on the command line

```
In [251]: # make the files executable
!chmod +x mapper.py
!chmod +x reducer.py

# first try with just the reducer
!cat data/Consumer_test.csv | ~/w261/mapper.py > Outputs/testing.txt

# read the first couple lines of the testing file
!head Outputs/testing.txt
```

```
Disclosure      1
verification    1
of              1
debt            1
Disclosure      1
verification    1
of              1
debt            1
Deposits        1
and             1
```

```
In [252]: # now let's test the output of the mapper into
# the reducer
!cat data/Consumer_test.csv | ~/w261/mapper.py | ~/w261/reducer.py > Out
puts/testing.txt

# read the first couple lines of the testing file
!tail Outputs/testing.txt
```

```
health  26
Cont'd  26
club    26
owed    26
debt    34
on       74
Incorrect      74
information    74
report  77
credit  97
```

Make the files executable and run the hadoop with 2 reducers

```
In [253]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_test.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-1
!rm -r /home/cloudera/w261/Outputs/Out_3_2-1

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-1

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=2 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-combiner reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-1

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-1/* /home/cloudera/w261/Outputs/Out_3_2-1
```



```

16/06/04 08:32:41 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart
t.cloudera:8020/user/cloudera/w261/ProductPurchaseData_test.txt' to tra
sh at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/use
r/cloudera/w261/ProductPurchaseData_test.txt1465054361377
rm: `/user/cloudera/w261-output-3-2-1': No such file or directory
rm: cannot remove `/home/cloudera/w261/Outputs/Out_3_2-1': No such file
or directory
16/06/04 08:32:57 INFO Configuration.deprecation: mapred.map.tasks is d
eprecated. Instead, use mapreduce.job.maps
16/06/04 08:32:57 INFO Configuration.deprecation: mapred.reduce.tasks i
s deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/
streamjob7077778030319964414.jar tmpDir=null
16/06/04 08:32:59 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/127.0.0.1:8032
16/06/04 08:32:59 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/127.0.0.1:8032
16/06/04 08:33:02 INFO mapred.FileInputFormat: Total input paths to pro
cess : 1
16/06/04 08:33:02 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 08:33:03 INFO mapreduce.JobSubmitter: Submitting tokens for jo
b: job_1464634906532_0017
16/06/04 08:33:05 INFO impl.YarnClientImpl: Submitted application appli
cation_1464634906532_0017
16/06/04 08:33:05 INFO mapreduce.Job: The url to track the job: http://
quickstart.cloudera:8088/proxy/application_1464634906532_0017/
16/06/04 08:33:06 INFO mapreduce.Job: Running job: job_1464634906532_00
17
16/06/04 08:33:24 INFO mapreduce.Job: Job job_1464634906532_0017 runnin
g in uber mode : false
16/06/04 08:33:24 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 08:34:27 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 08:34:54 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 08:34:57 INFO mapreduce.Job: Job job_1464634906532_0017 comple
ted successfully
16/06/04 08:34:57 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=1448
        FILE: Number of bytes written=486673
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=61947
        HDFS: Number of bytes written=1093
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=153
65760
        Total time spent by all reduces in occupied slots (ms)=
6151808
        Total time spent by all map tasks (ms)=120045
        Total time spent by all reduce tasks (ms)=48061

```

Total vcore-seconds taken by all map tasks=120045
 Total vcore-seconds taken by all reduce tasks=48061
 Total megabyte-seconds taken by all map tasks=15365760
 Total megabyte-seconds taken by all reduce tasks=615180

8

Map-Reduce Framework

Map input records=249
 Map output records=1083
 Map output bytes=11212
 Map output materialized bytes=1847
 Input split bytes=240
 Combine input records=1083
 Combine output records=159
 Reduce input groups=159
 Reduce shuffle bytes=1847
 Reduce input records=159
 Reduce output records=102
 Spilled Records=318
 Shuffled Maps =4
 Failed Shuffles=0
 Merged Map outputs=4
 GC time elapsed (ms)=1311
 CPU time spent (ms)=5560
 Physical memory (bytes) snapshot=501768192
 Virtual memory (bytes) snapshot=2943021056
 Total committed heap usage (bytes)=190840832

Shuffle Errors

BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0

File Input Format Counters

Bytes Read=61707

File Output Format Counters

Bytes Written=1093

16/06/04 08:34:57 INFO streaming.StreamJob: Output directory: /user/clo
 udera/w261-output-3-2-1

Run the full data through Hadoop

```
In [288]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/Consumer_Complaints.csv /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-2-1
!rm -r /home/cloudera/w261/Outputs/Out_3_2-1

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_2-1

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.map.tasks=2 -D mapred.reduce.tasks=2 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-combiner reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-2-1

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-2-1/* /home/cloudera/w261/Outputs/Out_3_2-1
```



```

16/06/04 10:30:51 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/Consumer_Complaints.csv' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/Consumer_Complaints.csv1465061451380
16/06/04 10:31:00 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-2-1' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-2-1
16/06/04 10:31:03 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
16/06/04 10:31:03 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob4751754695751656259.jar tmpDir=null
16/06/04 10:31:05 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:31:06 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:31:07 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 10:31:07 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 10:31:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0022
16/06/04 10:31:08 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0022
16/06/04 10:31:08 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0022/
16/06/04 10:31:08 INFO mapreduce.Job: Running job: job_1464634906532_0022
16/06/04 10:31:18 INFO mapreduce.Job: Job job_1464634906532_0022 running in uber mode : false
16/06/04 10:31:18 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 10:31:47 INFO mapreduce.Job:  map 25% reduce 0%
16/06/04 10:31:48 INFO mapreduce.Job:  map 59% reduce 0%
16/06/04 10:31:49 INFO mapreduce.Job:  map 75% reduce 0%
16/06/04 10:31:50 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 10:32:05 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 10:32:06 INFO mapreduce.Job: Job job_1464634906532_0022 completed successfully
16/06/04 10:32:07 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=12127
        FILE: Number of bytes written=499029
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50946999
        HDFS: Number of bytes written=2477
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=768

```

6400


```

Total time spent by all reduces in occupied slots (ms)=
3129856
Total time spent by all map tasks (ms)=60050
Total time spent by all reduce tasks (ms)=24452
Total vcore-seconds taken by all map tasks=60050
Total vcore-seconds taken by all reduce tasks=24452
Total megabyte-seconds taken by all map tasks=7686400
Total megabyte-seconds taken by all reduce tasks=312985

```

6

Map-Reduce Framework

```

Map input records=312913
Map output records=978634
Map output bytes=10407527
Map output materialized bytes=5984
Input split bytes=252
Combine input records=978634
Combine output records=508
Reduce input groups=508
Reduce shuffle bytes=5984
Reduce input records=508
Reduce output records=182
Spilled Records=1369
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=1239
CPU time spent (ms)=9540
Physical memory (bytes) snapshot=473772032
Virtual memory (bytes) snapshot=2924826624
Total committed heap usage (bytes)=187170816

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```
Bytes Read=50946747
```

File Output Format Counters

```
Bytes Written=2477
```

```

16/06/04 10:32:07 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-2-1

```

Pull out the top 50

In [289]:

```
# pull the two files that we outputted from
# hadoop
file1 = "Outputs/Out_3_2-1/part-00000"
file2 = "Outputs/Out_3_2-1/part-00001"

# set the lists for holding the words and
# the counts
words1 = []
counts1 = []
words2 = []
counts2 = []

# get the words from file 1
with open(file1, "r") as myfile:
    for line in myfile.readlines():

        # split the line
        line = line.split("\t")

        # set the word and the count
        word = line[0].strip()
        count = int(line[1].strip())

        # append to the lists
        words1.append(word)
        counts1.append(count)

# get the words from file 2
with open(file2, "r") as myfile:
    for line in myfile.readlines():

        # split the line
        line = line.split("\t")

        # set the word and the count
        word = line[0].strip()
        count = int(line[1].strip())

        # append to the lists
        words2.append(word)
        counts2.append(count)

# create a final list to store the counts
# and words
words = []
counts = []
LENGTH = 50

# while we have less than 50 items in
# our list
while len(words) < LENGTH:

    # compare the counts and words at
    # the end of each list
    word1 = words1[-1]
    count1 = counts1[-1]
    word2 = words2[-1]
```

```
count2 = counts2[-1]

# if word1 is larger, append it to the
# final list
if count1 > count2:
    words.append(words1.pop())
    counts.append(counts1.pop())

# else if word is actually bigger, append
# it to the final list
elif count2 > count1:
    words.append(words2.pop())
    counts.append(counts2.pop())

# else if the two words have equal counts
else:

    # append both words to a new list and
    # sort that list
    sort = []
    sort.append(word1)
    sort.append(word2)
    sort.sort()

    # check which word is alphabetically
    # first and add that word to our
    # final list
    if word1 == sort[0]:
        words.append(words1.pop())
        counts.append(counts1.pop())
    else:
        words.append(words2.pop())
        counts.append(counts2.pop())

# print out the list of the top 50 words
for index,word in enumerate(words):
    info = word + "\t" + str(counts[index])
    print info
```

"Loan	107254
modification	70487
servicing	36767
credit	36126
report	30546
Incorrect	29069
information	29069
on	29069
or	22533
debt	17966
and	16448
"Account	16205
opening	16205
Credit	14768
club	12545
health	12545
/	12386
not	12353
loan	12237
collect	11848
attempts	11848
owed	11848
Cont'd	11848
of	10885
my	10731
Deposits	10555
withdrawals	10555
Problems	9484
"Application	8625
to	8401
Billing	8158
Other	7886
disputes	6938
Communication	6920
tactics	6920
reporting	6559
lease	6337
the	6248
by	5663
caused	5663
being	5663
funds	5663
low	5663
process	5505
verification	5214
Disclosure	5214
Managing	5006
company's	4858
investigation	4858
card	4405

HW 3.3 Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers. For this homework use the online browsing behavior dataset located at:

<https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0>

[\(https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0\)](https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0)

Each line in this dataset represents a browsing session of a customer. On each line, each string of 8 characters represents the id of an item browsed during that session. The items are separated by spaces.

Here are the first few lines of the ProductPurchaseData:

FRO11987 ELE17451 ELE89019 SNA90258 GRO99222

GRO99222 GRO12298 FRO12685 ELE91550 SNA11465 ELE26917 ELE52966 FRO90334 SNA30755 ELE17451 FRO84225 SNA80192

ELE17451 GRO73461 DAI22896 SNA99873 FRO86643

ELE17451 ELE37798 FRO86643 GRO56989 ELE23393 SNA11465

ELE17451 SNA69641 FRO86643 FRO78087 SNA11465 GRO39357 ELE28573 ELE11375 DAI54444

Do some exploratory data analysis of this dataset guided by the following questions:

Using a single reducer: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

Preprocess the data

We want to preprocess the file to add a unique identifier for each basket. This allows the products in each basket to be sorted but still retain their association with the original basket. This is a step similar to the one described when dealing with Microsoft.com's log files in the Async.

```
In [290]: def Addkey(filename):  
    """Adds a unique key to every row in a  
    data file"""  
  
    # open the original file  
    with open(filename,"r") as myfile:  
  
        # create a new filename  
        newfilename = filename + "_mod"  
  
        # create a new file for the found  
        with open(newfilename,"w") as mynewfile:  
  
            # initialize a record counter to assign each  
            # basket it's own unique identifier  
            count = 0  
  
            # loop through all the lines  
            for line in myfile.readlines():  
  
                # create what we will be writing  
                info = str(count) + "..\t.." + line  
  
                # write to the new file  
                mynewfile.write(info)  
  
                # increment the count  
                count = count + 1
```

```
In [291]: # add a key to both my test and full data  
Addkey("data/ProductPurchaseData_test.txt")  
Addkey("data/ProductPurchaseData.txt")  
print "Keys successfully added"
```

Keys successfully added

Mapper function

The mapper function takes the browsing session from a website and outputs each product visited and the number of products visited.

```
In [293]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.3

# import the system library to read from
# the input
import sys

# loop through each line
for line in sys.stdin:

    # split by the divider
    line = line.split("..\\t..")

    # set the basket number and the items
    # in each basket
    basket_id = int(line[0])
    products = line[1].strip().split()

    # count the number of products in
    # the basket
    number = len(products)

    # print out the key-value pair as a tab
    # delimited list of word and
    for product in products:
        info = str(basket_id) + "\\t" + \\
        product.strip() + "," + str(number)
        print info
```

Overwriting mapper.py

Reducer function

The reducer function takes the output from the mapper and outputs:

- a list of products sorted by frequency of order, along with relative frequency
- number of unique products
- contents of the largest basket

In [294]:

```
%%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.3

# import the system library to read from
# the input
# import the operator library to help sort
# the dictionary
import sys
import operator

# create a dictionary to store the products
# and their respective counts
products = {}

# create an array to hold the products in the
# largest basket, also add a marker to help
# us keep track of the items we add to the basket
basket = []
max_basket_id = None

# keep a running count of all products
all_prods = 0

# loop through each line
for line in sys.stdin:

    # split the line by the tab and set the
    # basket id
    line = line.split("\t")
    _basket_id = line[0]

    # set the product the next item
    # and the number of items as the one
    # after that
    _product = line[1].split(",")[0]
    _basket_total = int(line[1].split(",")[1])

    # check to see if the product already
    # exists in the dictionary
    if _product not in products:

        # if it's not, add it, and initialize
        # it with a count of 0
        products[_product] = 0

    # increment the counts
    products[_product] = products[_product] + 1
    all_prods = all_prods + 1

    # compare the length of the current basket
    # with the length of the basket currently stored
    if _basket_total > len(basket):

        # set how many items we want to add to
```

```
# this basket
max_basket_id = _basket_id

# if we need to add items to the basket
if _basket_id == max_basket_id:

    # add the item to the basket
    basket.append(_product)

# print the sorted list of products
for product in sorted(products.items(),\
                        key=lambda k: (-k[1], k[0])):
    freq = float(product[1])/float(all_prods)
    print product[0],"\t",product[1],"\t",freq

# print a dividing line
print "*~*~*~*~*"

# print the length of the largest basket
print "The number of items in the largest basket \
is", len(basket)

# print each item in the basket
info = ""
for _item in basket:
    info = info + str(_item) + " "
print info

# print a dividing line
print "*~*~*~*~*"

# print the number of unique products
print "Unique products:", len(products.keys())
```

Overwriting reducer.py

Make the files executable and run the test data hadoop

```
In [62]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/ProductPurchaseData_te
st.txt_mod /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-3
!rm -r /home/cloudera/w261/Outputs/Out_3_3

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_3

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-stre
aming-mr1.jar \
-D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-3

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-3/part-00000 /home/c
loudera/w261/Outputs/Out_3_3
```



```

16/05/30 15:43:48 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/ProductPurchase_mod.txt' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/ProductPurchase_mod.txt
16/05/30 15:43:56 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-3' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-3
16/05/30 15:43:59 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob6949182719773429115.jar tmpDir=null
16/05/30 15:44:00 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 15:44:01 INFO client.RMPProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/05/30 15:44:02 INFO mapred.FileInputFormat: Total input paths to process : 1
16/05/30 15:44:03 INFO mapreduce.JobSubmitter: number of splits:2
16/05/30 15:44:03 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0007
16/05/30 15:44:04 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0007
16/05/30 15:44:04 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0007/
16/05/30 15:44:04 INFO mapreduce.Job: Running job: job_1464634906532_0007
16/05/30 15:44:27 INFO mapreduce.Job: Job job_1464634906532_0007 running in uber mode : false
16/05/30 15:44:27 INFO mapreduce.Job:  map 0% reduce 0%
16/05/30 15:45:21 INFO mapreduce.Job:  map 100% reduce 0%
16/05/30 15:45:31 INFO mapreduce.Job:  map 100% reduce 100%
16/05/30 15:45:31 INFO mapreduce.Job: Job job_1464634906532_0007 completed successfully
16/05/30 15:45:31 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=9505
        FILE: Number of bytes written=380822
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=18582
        HDFS: Number of bytes written=13913
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=11746944
        Total time spent by all reduces in occupied slots (ms)=956160
        Total time spent by all map tasks (ms)=91773
        Total time spent by all reduce tasks (ms)=7470
        Total vcore-seconds taken by all map tasks=91773

```

```
Total vcore-seconds taken by all reduce tasks=7470
Total megabyte-seconds taken by all map tasks=11746944
Total megabyte-seconds taken by all reduce tasks=956160
Map-Reduce Framework
  Map input records=105
  Map output records=1265
  Map output bytes=18709
  Map output materialized bytes=9749
  Input split bytes=252
  Combine input records=0
  Combine output records=0
  Reduce input groups=105
  Reduce shuffle bytes=9749
  Reduce input records=1265
  Reduce output records=447
  Spilled Records=2530
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=3090
  CPU time spent (ms)=3960
  Physical memory (bytes) snapshot=361803776
  Virtual memory (bytes) snapshot=2188435456
  Total committed heap usage (bytes)=142606336
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=18330
File Output Format Counters
  Bytes Written=13913
16/05/30 15:45:31 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-3
```

Run the full data through Hadoop

```
In [295]: # first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/ProductPurchaseData.txt_mod /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-3
!rm -r /home/cloudera/w261/Outputs/Out_3_3

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_3

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-3

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-3/part-00000 /home/cloudera/w261/Outputs/Out_3_3
```



```

16/06/04 10:45:35 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/ProductPurchaseData.txt_mod' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/ProductPurchaseData.txt_mod
16/06/04 10:45:44 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-3' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-31465062344300
16/06/04 10:45:47 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob8907208976098100491.jar tmpDir=null
16/06/04 10:45:48 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:45:49 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/127.0.0.1:8032
16/06/04 10:45:50 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 10:45:50 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 10:45:51 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1464634906532_0024
16/06/04 10:45:51 INFO impl.YarnClientImpl: Submitted application application_1464634906532_0024
16/06/04 10:45:51 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1464634906532_0024/
16/06/04 10:45:51 INFO mapreduce.Job: Running job: job_1464634906532_0024
16/06/04 10:46:03 INFO mapreduce.Job: Job job_1464634906532_0024 running in uber mode : false
16/06/04 10:46:03 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 10:46:34 INFO mapreduce.Job:  map 67% reduce 0%
16/06/04 10:46:36 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 10:46:50 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 10:46:51 INFO mapreduce.Job: Job job_1464634906532_0024 completed successfully
16/06/04 10:46:51 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=3065209
        FILE: Number of bytes written=6467348
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3780013
        HDFS: Number of bytes written=394555
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=7957248
        Total time spent by all reduces in occupied slots (ms)=1406592
        Total time spent by all map tasks (ms)=62166
        Total time spent by all reduce tasks (ms)=10989

```

```

Total vcore-seconds taken by all map tasks=62166
Total vcore-seconds taken by all reduce tasks=10989
Total megabyte-seconds taken by all map tasks=7957248
Total megabyte-seconds taken by all reduce tasks=140659

```

2

Map-Reduce Framework

```

Map input records=31101
Map output records=380824
Map output bytes=6631793
Map output materialized bytes=3040571
Input split bytes=260
Combine input records=0
Combine output records=0
Reduce input groups=31101
Reduce shuffle bytes=3040571
Reduce input records=380824
Reduce output records=12597
Spilled Records=761648
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=735
CPU time spent (ms)=8560
Physical memory (bytes) snapshot=376811520
Virtual memory (bytes) snapshot=2195718144
Total committed heap usage (bytes)=143654912

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```
Bytes Read=3779753
```

File Output Format Counters

```
Bytes Written=394555
```

```

16/06/04 10:46:51 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-3

```

Handle the output data

The output data has the answer to multiple questions. We split the file into multiple files, one for each part of the question.

```

In [296]: # set the divider information
divider = "*****"
divider_count = 0

# initialize variables to hold the information
# for each future file
products = []
largest = []
uniques = []

# open the output file
with open("Outputs/Out_3_3/part-00000","r") \
as myfile:

    # loop through every line
    for line in myfile.readlines():

        # check to see if we've reached a divider
        if line.strip() == divider:
            divider_count = divider_count + 1
        else:

            # add the line to the appropriate file
            # based on the divider count
            if divider_count == 0:
                products.append(line)
            elif divider_count == 1:
                largest.append(line)
            else:
                uniques.append(line)

# write to each of the new files
with open("Outputs/Out_3_3/products","w") as myfile:
    for product in products:
        myfile.write(product)
with open("Outputs/Out_3_3/largest","w") as myfile:
    for line in largest:
        myfile.write(line)
with open("Outputs/Out_3_3/uniques","w") as myfile:
    for line in uniques:
        myfile.write(line)

```

Number of unique products

```
In [297]: !cat Outputs/Out_3_3/uniques
```

```
Unique products: 12592
```

Largest basket

```
In [298]: !cat Outputs/Out_3_3/largest
```

```
The number of items in the largest basket is 71
```

```
ELE89019 FRO11987 ELE17451 GRO99222 SNA90258 FRO12685 GRO12298 GRO99222  
SNA80192 FRO84225 ELE17451 SNA30755 FRO90334 ELE52966 ELE26917 SNA1146  
5 ELE91550 DAI92253 DAI93692 SNA55952 GRO12935 GRO48282 DAI87514 SNA947  
81 SNA17715 ELE82555 GRO36567 SNA47306 GRO99222 DAI22896 GRO73461 ELE17  
451 DAI22177 ELE26917 FRO82427 ELE24180 GRO32086 ELE36890 ELE56095 SNA5  
5762 SNA80324 SNA72462 ELE87243 DAI67621 ELE20847 SNA61380 FRO15030 DAI  
75645 GRO99863 DAI38969 DAI62779 ELE56788 ELE81346 GRO94758 ELE49801 GR  
O68067 SNA47306 ELE59028 GRO69543 DAI53152 FRO84460 GRO81087 GRO61133 D  
AI85309 DAI84511 DAI54320 FRO37721 GRO46627 SNA96364 ELE35632 DAI67347
```

Top 50 products with their frequencies and relative frequencies

```
In [299]: print "Product\t\tCount\tRelative Frequency"
!head -50 Outputs/Out_3_3/products
```

Product	Count	Relative Frequency
DAI62779	6667	0.0175067747831
FRO40251	3881	0.010191059387
ELE17451	3875	0.0101753040775
GRO73461	3602	0.00945843749344
SNA80324	3044	0.00799319370628
ELE32164	2851	0.0074863979161
DAI75645	2736	0.00718442114993
SNA45677	2455	0.0064465474865
FRO31317	2330	0.0061183118711
DAI85309	2293	0.00602115412894
ELE26917	2292	0.00601852824402
FRO80039	2233	0.00586360103355
GRO21487	2115	0.00555374661261
SNA99873	2083	0.00546971829507
GRO59710	2004	0.00526227338613
GRO71621	1920	0.00504169905258
FRO85978	1918	0.00503644728273
GRO30386	1840	0.00483162825872
ELE74009	1816	0.00476860702057
GRO56726	1784	0.00468457870302
DAI63921	1773	0.00465569396887
GRO46854	1756	0.00461105392517
ELE66600	1713	0.00449814087347
DAI83733	1712	0.00449551498855
FRO32293	1702	0.00446925613932
ELE66810	1697	0.0044561267147
SNA55762	1646	0.00432220658362
DAI22177	1627	0.00427231477008
FRO78087	1531	0.00402022981745
ELE99737	1516	0.0039808415436
ELE34057	1489	0.00390994265067
GRO94758	1489	0.00390994265067
FRO35904	1436	0.00377077074974
FRO53271	1420	0.00372875659097
SNA93860	1407	0.00369462008697
SNA90094	1390	0.00364998004327
GRO38814	1352	0.00355019641619
ELE56788	1345	0.00353181522173
GRO61133	1321	0.00346879398357
DAI88807	1316	0.00345566455896
ELE74482	1316	0.00345566455896
ELE59935	1311	0.00344253513434
SNA96271	1295	0.00340052097557
DAI43223	1290	0.00338739155095
ELE91337	1289	0.00338476566603
GRO15017	1275	0.0033480032771
DAI31081	1261	0.00331124088818
GRO81087	1220	0.00320357960633
DAI22896	1219	0.0032009537214
GRO85051	1214	0.00318782429679

HW3.4. (Computationally prohibitive but then again Hadoop can handle this) Pairs

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a map-reduce program to find products which are frequently browsed together. Fix the support count (cooccurrence count) to $s = 100$ (i.e. product pairs need to occur together at least 100 times to be considered frequent) and find pairs of items (sometimes referred to itemsets of size 2 in association rule mining) that have a support count of 100 or more.

List the top 50 product pairs with corresponding support count (aka frequency), and relative frequency or support (number of records where they occur, the number of records where they occur/the number of baskets in the dataset) in decreasing order of support for frequent ($100 > \text{count}$) itemsets of size 2.

Use the Pairs pattern (lecture 3) to extract these frequent itemsets of size 2. Free free to use combiners if they bring value. Instrument your code with counters for count the number of times your mapper, combiner and reducers are called.

Please output records of the following form for the top 50 pairs (itemsets of size 2):
item1, item2, support count, support

Fix the ordering of the pairs lexicographically (left to right), and break ties in support (between pairs, if any exist) by taking the first ones in lexicographically increasing order.

Report the compute time for the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers).

Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts.

Mapper function

This mapper function takes the items in a basket and outputs pairs for each co-occurrence for every term. For example, the line a,b,c would output: ab 1, ac 1, ba 1, bc 1, ca 1, cb 1. This is the pairs method.

```
In [3]: %%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.4

# import the system library to read from
# the input
import sys

# initialize counter to keep track of the number
# of baskets
baskets = 0

# write to standard error to keep track of how often
# we call this function
sys.stderr.write("reporter:counter:MyJob,Mapper,1\n")

# loop through each line
for line in sys.stdin:

    # increment the counter
    baskets = baskets + 1

    # split by spaces
    products = line.split()


    # loop through each product
    for product in products:

        # loop through every other product,
        # ignoring the the product itself
        # because we wouldn't want
        # to make a product recommendation for the
        # same product we're already visiting
        for pair in products:

            # if this is another product
            if product.strip() != pair.strip():

                # create the string to print
                info = product.strip() + "," + \
                    pair.strip() + "\t1"
                print info

# print the counter as the last line
info = "Records:" + "\t" + str(baskets)
print info
```

Overwriting mapper.py

Reducer function

This function takes the output from the mappers and outputs in a sorted list for each co-occurring pair:

- item1
- item2
- support count (number of times this pair occurs)
- support (portion of baskets this pair occurs in)

In [5]:

```
%%writefile reducer.py
#!/home/cloudera/anaconda2/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.4

# import the system library to read from
# the input
# import the numpy library to help us sort
# before output
import sys
import numpy as np

# create a dictionary to store the product pairs
# and their respective counts
products = {}

# keep a running count of how many records we are
# dealing with
records = 0

# write a line to the standard error to keep
# track of how often this function is called
sys.stderr.write("reporter:counter:MyJob,Reducer,1\n")

# loop through each line
for line in sys.stdin:

    # split the line by the tab
    line = line.split("\t")

    # set the pair and count values
    pair = line[0].strip()
    count = int(line[1])

    # if its the count variable let's update
    # our counts
    if pair == "Records:":
        records = records + count

    else:
        # check to see if the pair already
        # exists in the dictionary
        if pair not in products:

            # if it's not, add it, and initialize
            # it with a count of 0
            products[pair] = 0

        # increment the counts
        products[pair] = products[pair] + 1

# create a set of array to store those pairs that
# frequently occur together
FREQUENT = 100
item1s = []
item2s = []
```

```
counts = []
frequs = []

# loop through all the pairs in the dictionary
for pair in products.keys():

    # if the pair occurs more than 100 times
    if int(products[pair]) > FREQUENT:

        # split the items and store them
        item1 = pair.split(',')[0]
        item2 = pair.split(',')[1]

        # get the count and relative
        # frequency
        count = products[pair]
        frequ = float(count)/float(records)

        # append all the values to our arrays
        item1s.append(item1)
        item2s.append(item2)
        counts.append(count)
        frequs.append(frequ)

# convert our arrays to numpy arrays
item1s = np.array(item1s)
item2s = np.array(item2s)
counts = np.array(counts)
frequs = np.array(frequs)

# get the indices for the sorted list
indexs = np.argsort(frequs)[::-1]

# print the elements to the output
print "Item1\t\tItem2\t\tCount\tSupport"
for index in indexs:

    # gather all the information
    item1 = item1s[index]
    item2 = item2s[index]
    count = counts[index]
    frequ = frequs[index]

    # create the line to print out
    line = str(item1) + "\t" + str(item2) + \
"\t" + str(count) + "\t" + str(frequ)

    # print the line out
    print line
```

Overwriting reducer.py

Make the files executable and run the test data in hadoop

```
In [1]: # write the time to a file
!echo $(date) > Outputs/Time_3_4

# first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/ProductPurchaseData_test.txt /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-4
!rm -r /home/cloudera/w261/Outputs/Out_3_4

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_4

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-4

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-4/part-00000 /home/cloudera/w261/Outputs/Out_3_4

# write the ending time to a file
!echo $(date) >> Outputs/Time_3_4
```



```

16/06/04 12:46:54 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/ProductPurchaseData_test.txt' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/ProductPurchaseData_test.txt1465069614349
16/06/04 12:47:02 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-4' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-41465069622054
16/06/04 12:47:04 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob2656552094061797701.jar tmpDir=null
16/06/04 12:47:05 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/06/04 12:47:05 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/06/04 12:47:06 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 12:47:06 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 12:47:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1465069396819_0001
16/06/04 12:47:07 INFO impl.YarnClientImpl: Submitted application application_1465069396819_0001
16/06/04 12:47:08 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1465069396819_0001/
16/06/04 12:47:08 INFO mapreduce.Job: Running job: job_1465069396819_0001
16/06/04 12:47:17 INFO mapreduce.Job: Job job_1465069396819_0001 running in uber mode : false
16/06/04 12:47:17 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 12:47:29 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 12:47:36 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 12:47:37 INFO mapreduce.Job: Job job_1465069396819_0001 completed successfully
16/06/04 12:47:37 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=36092
        FILE: Number of bytes written=434757
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=8046
        HDFS: Number of bytes written=28
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=2490496
        Total time spent by all reduces in occupied slots (ms)=654848
        Total time spent by all map tasks (ms)=19457
        Total time spent by all reduce tasks (ms)=5116
        Total vcore-seconds taken by all map tasks=19457

```

```

Total vcore-seconds taken by all reduce tasks=5116
Total megabyte-seconds taken by all map tasks=2490496
Total megabyte-seconds taken by all reduce tasks=654848
Map-Reduce Framework
  Map input records=50
  Map output records=7026
  Map output bytes=140504
  Map output materialized bytes=37094
  Input split bytes=262
  Combine input records=0
  Combine output records=0
  Reduce input groups=5207
  Reduce shuffle bytes=37094
  Reduce input records=7026
  Reduce output records=1
  Spilled Records=14052
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=180
  CPU time spent (ms)=3690
  Physical memory (bytes) snapshot=428154880
  Virtual memory (bytes) snapshot=2210664448
  Total committed heap usage (bytes)=142082048
MyJob
  Mapper=2
  Reducer=1
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=7784
File Output Format Counters
  Bytes Written=28
16/06/04 12:47:37 INFO streaming.StreamJob: Output directory: /user/clo
udera/w261-output-3-4

```

Run the whole data through Hadoop


```
In [6]: # write the time to a file
!echo $(date) > Outputs/Time_3_4

# first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/ProductPurchaseData.txt /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-4
!rm -r /home/cloudera/w261/Outputs/Out_3_4

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_4

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-4

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-4/part-00000 /home/cloudera/w261/Outputs/Out_3_4

# write the ending time to a file
!echo $(date) >> Outputs/Time_3_4
```



```

16/06/04 12:54:34 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstar
t.cloudera:8020/user/cloudera/w261/ProductPurchaseData_test.txt' to tra
sh at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/use
r/cloudera/w261/ProductPurchaseData_test.txt1465070074023
16/06/04 12:54:40 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstar
t.cloudera:8020/user/cloudera/w261-output-3-4' to trash at: hdfs://quic
kstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-ou
tput-3-41465070080269
16/06/04 12:54:42 INFO Configuration.deprecation: mapred.reduce.tasks i
s deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/
streamjob9078575781275816040.jar tmpDir=null
16/06/04 12:54:43 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/10.0.2.15:8032
16/06/04 12:54:43 INFO client.RMProxy: Connecting to ResourceManager at
quickstart.cloudera/10.0.2.15:8032
16/06/04 12:54:45 INFO mapred.FileInputFormat: Total input paths to pro
cess : 1
16/06/04 12:54:45 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 12:54:45 INFO mapreduce.JobSubmitter: Submitting tokens for jo
b: job_1465069396819_0002
16/06/04 12:54:45 INFO impl.YarnClientImpl: Submitted application appli
cation_1465069396819_0002
16/06/04 12:54:45 INFO mapreduce.Job: The url to track the job: http://
quickstart.cloudera:8088/proxy/application_1465069396819_0002/
16/06/04 12:54:45 INFO mapreduce.Job: Running job: job_1465069396819_00
02
16/06/04 12:54:53 INFO mapreduce.Job: Job job_1465069396819_0002 runnin
g in uber mode : false
16/06/04 12:54:53 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 12:55:07 INFO mapreduce.Job:  map 51% reduce 0%
16/06/04 12:55:10 INFO mapreduce.Job:  map 59% reduce 0%
16/06/04 12:55:11 INFO mapreduce.Job:  map 77% reduce 0%
16/06/04 12:55:12 INFO mapreduce.Job:  map 83% reduce 0%
16/06/04 12:55:13 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 12:55:25 INFO mapreduce.Job:  map 100% reduce 83%
16/06/04 12:55:28 INFO mapreduce.Job:  map 100% reduce 93%
16/06/04 12:55:31 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 12:55:32 INFO mapreduce.Job: Job job_1465069396819_0002 comple
ted successfully
16/06/04 12:55:32 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=43244633
        FILE: Number of bytes written=62302003
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3498983
        HDFS: Number of bytes written=101764
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=442

```

8288

Total time spent by all reduces in occupied slots (ms)=

2150144

Total time spent by all map tasks (ms)=34596

Total time spent by all reduce tasks (ms)=16798

Total vcore-seconds taken by all map tasks=34596

Total vcore-seconds taken by all reduce tasks=16798

Total megabyte-seconds taken by all map tasks=4428288

Total megabyte-seconds taken by all reduce tasks=215014

4

Map-Reduce Framework

Map input records=31101

Map output records=5068110

Map output bytes=101362190

Map output materialized bytes=18550851

Input split bytes=252

Combine input records=0

Combine output records=0

Reduce input groups=1754191

Reduce shuffle bytes=18550851

Reduce input records=5068110

Reduce output records=2623

Spilled Records=15204330

Shuffled Maps =2

Failed Shuffles=0

Merged Map outputs=2

GC time elapsed (ms)=200

CPU time spent (ms)=31270

Physical memory (bytes) snapshot=431661056

Virtual memory (bytes) snapshot=2203160576

Total committed heap usage (bytes)=144179200

MyJob

Mapper=2

Reducer=1

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Input Format Counters

Bytes Read=3498731

File Output Format Counters

Bytes Written=101764

16/06/04 12:55:32 INFO streaming.StreamJob: Output directory: /user/clo
 udera/w261-output-3-4

Show the output

```
In [7]: !head -50 Outputs/Out_3_4/part-00000
```

Item1	Item2	Count	Support
ELE17451	DAI62779	1592	0.0511880646925
DAI62779	ELE17451	1592	0.0511880646925
FRO40251	SNA80324	1412	0.0454004694383
SNA80324	FRO40251	1412	0.0454004694383
FRO40251	DAI75645	1254	0.0403202469374
DAI75645	FRO40251	1254	0.0403202469374
FRO40251	GRO85051	1213	0.0390019613517
GRO85051	FRO40251	1213	0.0390019613517
GRO73461	DAI62779	1139	0.0366226166361
DAI62779	GRO73461	1139	0.0366226166361
DAI75645	SNA80324	1130	0.0363332368734
SNA80324	DAI75645	1130	0.0363332368734
FRO40251	DAI62779	1070	0.0344040384554
DAI62779	FRO40251	1070	0.0344040384554
DAI62779	SNA80324	923	0.0296775023311
SNA80324	DAI62779	923	0.0296775023311
DAI85309	DAI62779	918	0.0295167357963
DAI62779	DAI85309	918	0.0295167357963
GRO59710	ELE32164	911	0.0292916626475
ELE32164	GRO59710	911	0.0292916626475
GRO73461	FRO40251	882	0.0283592167454
DAI75645	DAI62779	882	0.0283592167454
DAI62779	DAI75645	882	0.0283592167454
FRO40251	GRO73461	882	0.0283592167454
DAI62779	ELE92920	877	0.0281984502106
ELE92920	DAI62779	877	0.0281984502106
FRO92469	FRO40251	835	0.026848011318
FRO40251	FRO92469	835	0.026848011318
ELE32164	DAI62779	832	0.0267515513971
DAI62779	ELE32164	832	0.0267515513971
GRO73461	DAI75645	712	0.0228931545609
DAI75645	GRO73461	712	0.0228931545609
ELE32164	DAI43223	711	0.022861001254
DAI43223	ELE32164	711	0.022861001254
GRO30386	DAI62779	709	0.02279669464
DAI62779	GRO30386	709	0.02279669464
ELE17451	FRO40251	697	0.0224108549564
FRO40251	ELE17451	697	0.0224108549564
ELE99737	DAI85309	659	0.0211890292917
DAI85309	ELE99737	659	0.0211890292917
DAI62779	ELE26917	650	0.020899649529
ELE26917	DAI62779	650	0.020899649529
GRO21487	GRO73461	631	0.0202887366966
GRO73461	GRO21487	631	0.0202887366966
SNA45677	DAI62779	604	0.0194205974084
DAI62779	SNA45677	604	0.0194205974084
ELE17451	SNA80324	597	0.0191955242597
SNA80324	ELE17451	597	0.0191955242597
DAI62779	GRO71621	595	0.0191312176457

Calculate the time the program took to run

```
In [8]: import datetime

def CalcTime(filepath):
    """A function that takes a simply file with
    only two lines, a starting and an ending time
    and returns the time elapsed between the two
    times"""

    # open the file
    with open(filepath,"r") as myfile:

        # set a starting and ending time
        start = None
        end = None

        # read in the lines and store them
        start = myfile.readline().strip()
        end = myfile.readline().strip()

        # set the time format
        time_format = "%a %b %d %H:%M:%S %Z %Y"

        # convert it to a time
        start = datetime.datetime.strptime(\
            start,\
            time_format)

        end = datetime.datetime.strptime(\
            end,\
            time_format)

        # calculate the elapsed time
        elapsed = end - start

        # return the elapsed time
        return elapsed
```

```
In [9]: # calculate the elapsed time
elapsed_3_4 = CalcTime("Outputs/Time_3_4")
print "The time it took to run the program \
\nis",elapsed_3_4
```

```
The time it took to run the program
is 0:01:05
```

Counters for mappers and reducers

We used a total of 3 counters, 2 for mappers and 1 for reducers.

MyJob	Name	Map	Reduce	Total
	Mapper	2	0	2
	Reducer	0	1	1

HW 3.5

Repeat 3.4 using the stripes design pattern for finding cooccurring pairs.

Report the compute times for stripes job versus the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers)

Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts. Discuss the differences in these counts between the Pairs and Stripes jobs.

Mapper function

The mapper function takes as input all items in a basket, each basket separated by a new line character. It outputs a dictionary key that holds the counts for each product and the different products it is paired with. For example, consider reading in the following list: a,b,c ; a,b ; c,d, ; a,b,e. Our output for product a will be: a (b:3, c:1, e:1).

In [1]:


```
%%writefile mapper.py
#!/usr/bin/python
## mapper.py
## Author: Alex Smith
## Description: mapper code for HW3.5

# import the system library to read from
# the input
import sys

# initialize counter to keep track of the number
# of baskets
baskets = 0

# initialize a dictionary to store each product
# and the counts for its paired products
prods = {}

# write to standard error to keep track of how often
# we call this function
sys.stderr.write("reporter:counter:MyJob,Mapper,1\n")

# loop through each line
for line in sys.stdin:

    # increment the counter
    baskets = baskets + 1

    # split by spaces
    products = line.split()

    # loop through each product
    for product in products:

        # strip away any excess
        product = product.strip()

        # if the product is not in already added
        # to the dictionary, add it
        if product not in prods.keys():
            prods[product]={}

        # loop through every other product,
        # ignoring the the product itself
        # because we wouldn't want
        # to make a product recomendation for the
        # same product we're already visiting
        for pair in products:

            # strip away any excess
            pair = pair.strip()

            # if this is another product
            if product != pair:

                # check to see if we already have
                # this pair in this product's
```

```
# dictionary
if pair not in prods[product].keys():
    prods[product][pair] = 0

# increment the count for the pairs
# we've seen with this specific product
prods[product][pair] = \
prods[product][pair] + 1

# print the dictionary for each product
for product in prods.keys():
    info = product + "\t" + str(prods[product])
    print info

# print the counter as the last line
info = "Records:" + "\t" + str(baskets)
print info
```

Overwriting mapper.py

```
In [2]: # test our mapper in the command line
!chmod +x mapper.py
!cat data/ProductPurchaseData_test.txt | ~/w261/mapper.py > Outputs/testing.txt
!head Outputs/testing.txt
```

```
reporter:counter:MyJob,Mapper,1
DAI50921      {'FRO78994': 1, 'GRO56989': 1, 'FRO32293': 1, 'FRO7121
3': 1, 'FRO18919': 1, 'SNA93730': 1, 'DAI36452': 1, 'ELE59935': 1, 'GRO
75578': 1, 'SNA14713': 1, 'GRO73461': 1, 'SNA58915': 1, 'GRO12935': 1,
'DAI35347': 1, 'SNA81153': 1, 'SNA99873': 1, 'SNA45677': 1, 'SNA8019
2': 1, 'ELE17451': 2, 'SNA89670': 1, 'DAI22896': 1, 'FRO41069': 1, 'DAI
62779': 1}
FRO70974      {'GRO30386': 1, 'DAI88808': 1, 'GRO36567': 1, 'SNA9155
4': 1, 'SNA80192': 1, 'FRO16142': 1, 'ELE17451': 1, 'ELE52446': 1, 'ELE
96863': 1, 'DAI22896': 1, 'SNA47306': 1, 'SNA90258': 1, 'FRO18919': 1,
'FRO41069': 1, 'GRO99222': 1, 'GRO49037': 1, 'FRO75418': 1, 'GRO7346
1': 1}
GRO94047      {'DAI22177': 1, 'FRO60023': 1, 'DAI59119': 1, 'FRO3131
7': 1, 'ELE89019': 1, 'ELE20196': 1, 'ELE17451': 1}
ELE13292      {'ELE12792': 1, 'DAI33885': 1, 'ELE27376': 1, 'DAI4435
5': 1, 'SNA56035': 1, 'ELE22574': 1, 'GRO71615': 1, 'ELE20166': 1, 'GRO
92942': 1}
ELE89019      {'GRO36567': 1, 'GRO56989': 1, 'GRO94047': 1, 'FRO3131
7': 1, 'ELE66067': 1, 'ELE14480': 1, 'FRO11987': 1, 'FRO36081': 1, 'DAI
93692': 1, 'ELE20196': 1, 'FRO60023': 1, 'SNA83730': 1, 'DAI22177': 1,
'DAI59119': 1, 'SNA85034': 1, 'ELE73246': 1, 'ELE59935': 2, 'DAI2253
4': 1, 'SNA63157': 1, 'ELE37798': 1, 'FRO92261': 1, 'SNA90258': 1, 'FRO
75418': 1, 'ELE96863': 1, 'DAI91290': 1, 'DAI35347': 1, 'GRO39070': 1,
'SNA55952': 1, 'DAI63921': 1, 'GRO82070': 1, 'SNA80192': 1, 'ELE1745
1': 5, 'FRO62970': 1, 'FRO79301': 1, 'DAI48891': 1, 'DAI62779': 1, 'GRO
99222': 2, 'GRO49037': 1, 'FRO78087': 2}
SNA11465      {'GRO56989': 2, 'GRO39369': 1, 'ELE11375': 1, 'DAI5444
4': 1, 'ELE91550': 1, 'FRO84225': 1, 'SNA69641': 1, 'FRO12685': 1, 'FRO
86643': 3, 'ELE23393': 1, 'FRO90334': 1, 'ELE52966': 1, 'ELE37798': 2,
'GRO39357': 2, 'ELE28573': 1, 'SNA30755': 1, 'FRO78087': 2, 'GRO1229
8': 1, 'GRO75578': 1, 'ELE26917': 1, 'SNA80192': 1, 'ELE17451': 4, 'GRO
99222': 1, 'DAI95741': 1}
GRO27756      {'SNA99873': 1, 'ELE91337': 1, 'DAI69239': 1, 'ELE5296
6': 1, 'GRO99222': 1, 'SNA14713': 1, 'SNA93641': 1}
FRO36081      {'ELE96863': 1, 'FRO92261': 1, 'GRO56989': 1, 'FRO7541
8': 1, 'SNA80192': 1, 'ELE17451': 1, 'ELE89019': 1, 'ELE73246': 1, 'ELE
66067': 1, 'DAI48891': 1, 'DAI35347': 1, 'GRO36567': 1, 'SNA55952': 1,
'DAI93692': 1, 'GRO49037': 1, 'SNA63157': 1, 'ELE37798': 1}
ELE20196      {'DAI22177': 1, 'DAI59119': 1, 'GRO94047': 1, 'FRO3131
7': 1, 'ELE89019': 1, 'FRO60023': 1, 'ELE17451': 1}
DAI84001      {'SNA45677': 1, 'GRO17442': 1, 'SNA77101': 1, 'SNA5891
5': 1, 'ELE91337': 1, 'DAI52318': 1, 'SNA30579': 1, 'ELE11468': 1, 'FRO
35353': 1, 'ELE20166': 1}
```

Reducer function

The reducer function takes the dictionaries outputted by the mapper and returns a sorted list for each co-occurring pair:

- item1
- item2
- support count (number of times this pair occurs)
- support (portion of baskets this pair occurs in)

In [7]:

```
%%writefile reducer.py
#!/usr/bin/python
## reducer.py
## Author: Alex Smith
## Description: reducer code for HW3.5

# import the system library to read from
# the input
# import the numpy library to help us sort
# before output
import sys
import numpy as np

# create a dictionary to store the product pairs
# and their respective counts
products = {}

# keep a running count of how many records we are
# dealing with
records = 0

# write a line to the standard error to keep
# track of how often this function is called
sys.stderr.write("reporter:counter:MyJob,Reducer,1\n")

# loop through each line
for line in sys.stdin:

    # split the line by the tab
    line = line.split("\t")

    # check to see if this is our records counter
    if line[0].strip() == "Records:":
        records = records + int(line[1])

    # otherwise we know its a pair we need to
    # process
    else:

        # set the first item
        item1 = line[0]

        # first let's convert it back into
        # a dictionary
        paired = eval(line[1])

        # loop through each of the paired words
        for pair in paired.keys():

            # set the count for the pair
            count = paired[pair]

            # create the pair
            pairs = item1 + "," + pair

            # check to see if it's already in
            # the dictionary
```

```
        if pairs not in products.keys():
            products[pairs] = 0

        # add to the count of the dictionary
        products[pairs] = products[pairs] + \
            count

# create a set of array to store those pairs that
# frequently occur together
FREQUENT = 100

# loop through all the pairs in the dictionary
for pair in products.keys():

    # if the pair occurs more than 100 times
    if int(products[pair]) > FREQUENT:

        # split the items and store them
        item1 = pair.split(',')[0]
        item2 = pair.split(',')[1]

        # get the count and relative
        # frequency
        count = products[pair]
        frequ = float(count)/float(records)

        # create the string to print
        info = item1 + "\t" + item2 + "\t" + \
            str(count) + "\t" + str(frequ)
        print info

# # convert our arrays to numpy arrays
# item1s = np.array(item1s)
# item2s = np.array(item2s)
# counts = np.array(counts)
# frequs = np.array(frequs)

# # get the indices for the sorted list
# indexs = np.argsort(frequs)[::-1]

# # print the elements to the output
# print "Item1\tItem2\tCount\tSupport"
# for index in indexs:

#     # gather all the information
#     item1 = item1s[index]
#     item2 = item2s[index]
#     count = counts[index]
#     frequ = frequs[index]

#     # create the line to print out
#     line = str(item1) + "\t" + str(item2) + \
#         "\t" + str(count) + "\t" + str(frequ)

#     # print the line out
#     print line
```

Overwriting reducer.py

```
In [54]: # test the reducer with the mapper's output
# slightly modify the reducer to lower the
# frequency to 10
!chmod +x reducer.py
!cat Outputs/testing.txt | ~/w261/reducer.py
```

```
reporter:counter:MyJob,Reducer,1
ELE17451      GRO99222      15      0.3
ELE17451      GRO73461      11      0.22
GRO73461      ELE17451      11      0.22
GRO99222      ELE17451      15      0.3
```

Make the files executable and run the test data in Hadoop


```
In [ ]: # write the current time to a file
!echo $(date) > Outputs/Time_3_5

# first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/ProductPurchaseData_test2.txt /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-5
!rm -r /home/cloudera/w261/Outputs/Out_3_5

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_5

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-5

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-5/part-00000 /home/cloudera/w261/Outputs/Out_3_5

# add the ending time to the file
!echo $(date) >> Outputs/Time_3_5
```

```
16/06/04 21:04:16 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/ProductPurchaseData.txt' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/ProductPurchaseData.txt1465099456855
16/06/04 21:04:29 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-5' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-51465099469400
16/06/04 21:04:34 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob5725157001918088713.jar tmpDir=null
16/06/04 21:04:35 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/06/04 21:04:36 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/06/04 21:04:38 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 21:04:38 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 21:04:39 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1465098020484_0003
16/06/04 21:04:39 INFO impl.YarnClientImpl: Submitted application application_1465098020484_0003
16/06/04 21:04:39 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1465098020484_0003/
16/06/04 21:04:39 INFO mapreduce.Job: Running job: job_1465098020484_0003
16/06/04 21:04:59 INFO mapreduce.Job: Job job_1465098020484_0003 running in uber mode : false
16/06/04 21:04:59 INFO mapreduce.Job: map 0% reduce 0%
```

Run the full data through Hadoop

```
In [5]: # write the current time to a file
!echo $(date) > Outputs/Time_3_5

# first let's clear our input directory to make
# sure that we're starting off with a clean slate
!hdfs dfs -rm -r /user/cloudera/w261/*

# modifies the permission to make the programs
# executable
!chmod +x ~/w261/mapper.py; chmod +x ~/w261/reducer.py

# put the input data into the hadoop cluster
!hdfs dfs -copyFromLocal /home/cloudera/w261/data/ProductPurchaseData_test.txt /user/cloudera/w261/

# make sure that we don't already have this output
!hdfs dfs -rm -r /user/cloudera/w261-output-3-5
!rm -r /home/cloudera/w261/Outputs/Out_3_5

# make the directory to store the output
!mkdir /home/cloudera/w261/Outputs/Out_3_5

# run the hadoop command
!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-mr1.jar \
-D mapred.reduce.tasks=1 \
-files /home/cloudera/w261/mapper.py,/home/cloudera/w261/reducer.py \
-mapper mapper.py \
-reducer reducer.py \
-input /user/cloudera/w261/* -output /user/cloudera/w261-output-3-5

# copy the output files to the local directory
!hdfs dfs -copyToLocal /user/cloudera/w261-output-3-5/part-00000 /home/cloudera/w261/Outputs/Out_3_5

# add the ending time to the file
!echo $(date) >> Outputs/Time_3_5
```




```

16/06/04 20:43:43 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261/ProductPurchaseData_test.txt' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261/ProductPurchaseData_test.txt1465098223712
16/06/04 20:43:52 INFO fs.TrashPolicyDefault: Moved: 'hdfs://quickstart.cloudera:8020/user/cloudera/w261-output-3-5' to trash at: hdfs://quickstart.cloudera:8020/user/cloudera/.Trash/Current/user/cloudera/w261-output-3-51465098232312
16/06/04 20:43:55 INFO Configuration.deprecation: mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamjob4761734599335409084.jar tmpDir=null
16/06/04 20:43:56 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/06/04 20:43:57 INFO client.RMProxy: Connecting to ResourceManager at quickstart.cloudera/10.0.2.15:8032
16/06/04 20:43:58 INFO mapred.FileInputFormat: Total input paths to process : 1
16/06/04 20:43:58 INFO mapreduce.JobSubmitter: number of splits:2
16/06/04 20:43:58 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1465098020484_0001
16/06/04 20:43:59 INFO impl.YarnClientImpl: Submitted application application_1465098020484_0001
16/06/04 20:43:59 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1465098020484_0001/
16/06/04 20:43:59 INFO mapreduce.Job: Running job: job_1465098020484_0001
16/06/04 20:44:09 INFO mapreduce.Job: Job job_1465098020484_0001 running in uber mode : false
16/06/04 20:44:09 INFO mapreduce.Job:  map 0% reduce 0%
16/06/04 20:44:22 INFO mapreduce.Job:  map 50% reduce 0%
16/06/04 20:44:23 INFO mapreduce.Job:  map 100% reduce 0%
16/06/04 20:44:33 INFO mapreduce.Job:  map 100% reduce 100%
16/06/04 20:44:34 INFO mapreduce.Job: Job job_1465098020484_0001 completed successfully
16/06/04 20:44:34 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=28067
        FILE: Number of bytes written=415836
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=8046
        HDFS: Number of bytes written=102
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=2945024
        Total time spent by all reduces in occupied slots (ms)=963840
        Total time spent by all map tasks (ms)=23008
        Total time spent by all reduce tasks (ms)=7530

```

Total vcore-seconds taken by all map tasks=23008
 Total vcore-seconds taken by all reduce tasks=7530
 Total megabyte-seconds taken by all map tasks=2945024
 Total megabyte-seconds taken by all reduce tasks=963840

Map-Reduce Framework

Map input records=50
 Map output records=268
 Map output bytes=86896
 Map output materialized bytes=26198
 Input split bytes=262
 Combine input records=0
 Combine output records=0
 Reduce input groups=219
 Reduce shuffle bytes=26198
 Reduce input records=268
 Reduce output records=4
 Spilled Records=536
 Shuffled Maps =2
 Failed Shuffles=0
 Merged Map outputs=2
 GC time elapsed (ms)=212
 CPU time spent (ms)=2640
 Physical memory (bytes) snapshot=416141312
 Virtual memory (bytes) snapshot=2192523264
 Total committed heap usage (bytes)=143130624

MyJob

Mapper=2
 Reducer=1

Shuffle Errors

BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0

File Input Format Counters

Bytes Read=7784

File Output Format Counters

Bytes Written=102

16/06/04 20:44:34 INFO streaming.StreamJob: Output directory: /user/cloudera/w261-output-3-5

Use of counters

We used a total of 3 counters, 2 for the mappers and 1 for the reducer. We would have had more instances of the reducer counter had we used a combiner. However, in this case, we elected to not use a combiner.

	Name ^	Map ↕	Reduce ↕	Total ↕
MyJob	Mapper	2	0	2
	Reducer	0	1	1

Calculate the elapsed time

```
In [231]: elapsed_3_5 = CalcTime("Outputs/Time_3_5")
print "The elapsed time for the stripes approach \
\nis",elapsed_3_5
```

The elapsed time for the stripes approach
is 0:01:13

```
In [233]: # let's compare the times
from prettytable import PrettyTable

# create a new table
pretty = PrettyTable(["Approach","Time"])
pretty.add_row(["Pairs",elapsed_3_4])
pretty.add_row(["Stripes",elapsed_3_5])

# show the table
print pretty
```

```
+-----+-----+
| Approach |    Time    |
+-----+-----+
|   Pairs   | 0:01:21    |
|  Stripes  | 0:01:13    |
+-----+-----+
```